

Protegis — Diligence, live. On merit.

VC diligence that takes analysts days, done live in ~3.2 seconds — on evidence, not who you know.

⚡ ~3.2S · SIGNAL → DECISION

Every claim checked

Against live GitHub, live web, live evidence — not cached data, not assumptions.

Three independent axes

Founder · Market · Idea-vs-Market — scored separately, never blended into a flattering composite.

One engine, no exceptions

The overstater, the honest unknown, the founder who never applied — same pipeline, same standard.

The Problem

The Claim Stack

A founder's deck arrives. "10,000 active users." "Ex-DeepMind." "Pilot with a Tier-1 bank." Some are true. Some are hope wearing a suit.

- GitHub tab — is the repo real?
- LinkedIn tab — did they actually work there?
- Crunchbase tab — any prior funding?
- Press tab — is the pilot announced anywhere?
- Repeat. For every claim. For every founder.



~2 days of tab-hopping per founder. Per round.

The Cost

By the time verification is done, the round has moved. Capital flows to who you know — not what's being built.

The overstater wins

Confident claims,
polished deck, warm
intro. Gets the meeting.

The builder loses

No network. Quietly
shipping something real.
Never gets read.

The founder with no network, quietly shipping something real, never gets read. That's the market failure Protegis fixes.

What Protegis Is

Drop in a deck — and optionally a résumé, GitHub, links. Protegis pulls out every claim and checks each one against live evidence in seconds.



What goes in

- Founder deck (PDF or link)
- Résumé (optional)
- GitHub profile (optional)
- Any additional links

What comes out

An investor-grade memo where every claim is marked one of three ways — each traced to the exact source that determined the verdict.

✓ Verified	❌ Contradicted
? Not enough evidence	

📌 One engine handles the overstater, the legit founder, the unknown with no track record, and founders who never applied.

Benefit 1 — Diligence in Seconds, Not Days

🕒 ~3.2S · SIGNAL → DECISION

A measured, instrumented number — not a slogan. The proof is in a single live example.

The Claim

"10,000 active users"

Stated confidently in the deck. Sounds credible. Sounds like traction. An analyst would spend hours verifying this across multiple sources.

Protegis pulls the real repository — live, on data it has never seen before — in the time it takes to read this card.

The Verdict

CONTRADICTED

14 stars

Not 10,000 users. The repo exists — but the numbers don't.

Dormant 7 months

Last commit: 7 months ago. No active development. No active users.

Receipt attached

github.com/... — exact source, exact timestamp, exact evidence.

A caught lie, with a receipt attached — in ~3.2 seconds.

Benefit 2 — Evidence You Can Actually Trust

A verdict is only as good as the reasoning behind it. Protegis makes that reasoning visible, checkable, and structurally resistant to inflation.



Trust Score

Every claim gets a status, a confidence level, and the exact evidence it rests on. Not a black box — a traceable judgment.



Independent Validator

A second pass re-checks verdicts against freshly fetched evidence. It can only ever *downgrade*, never upgrade. A checker that can't upgrade can't invent confidence.



Typed Audit Trail

Every verdict carries a full record — what evidence was used, what was checked, what got corrected. A judge sees the reasoning, not just the answer.

- ❏ Structural asymmetry is the key design choice: the validator can only downgrade. This removes the most common failure mode in AI scoring — confidence inflation.

Benefit 3 — Fair by Design

Three axes scored separately. Never blended. A brilliant founder in a dead market can't hide behind an average.



Founder

Track record, credibility, evidence of execution. Scored on what exists — with an honest confidence band when evidence is thin.



Market

Size, timing, competitive dynamics. A strong founder score can't rescue a dead market — they stay separate.



Idea-vs-Market

Does the idea fit the market being claimed? Scored independently so a mismatch is visible, not averaged away.

The hardest case, handled honestly

A founder with no track record still gets scored — with a wide, honest confidence band and a clear list of what would verify them further. No track record is not the same as a bad track record.

The honest number

64 ± 14

An honest "64, give or take 14" beats a confident number built on nothing. Wide band = honest. Narrow band = evidence-backed.

Benefit 4 — It Finds Founders Who Never Applied

Protegis doesn't just judge who shows up. Point it at your thesis and it scans GitHub live for founders quietly building what you're looking for.

The standard model

→ Founder submits deck

Only founders with networks, warm intros, or awareness of your fund ever reach you.

→ Analyst reviews

2 days of tab-hopping. Round moves. Opportunity missed.

→ Capital flows to who you know

The builder with no network never gets read.

The Protegis model

→ Point at your thesis

Define what you're looking for. Protegis scans GitHub live for founders building it — whether they know you exist or not.

→ Surface the unknown builder

Founders discovered this way are scored on the *same* engine as everyone else. No separate, lesser path.

→ Same engine. No exceptions.


Discovery doesn't lower the bar. It removes the gatekeeping before the bar.



Same engine. No separate path. The founder who never applied gets the same diligence as the one with a warm intro.

Benefit 5 — It Learns Where Great Founders Come From

Every outcome feeds back into the system. Protegis doesn't just score — it gets smarter about where to look next.



Ranked by quality, not volume

Shrinkage-adjusted so one lucky founder can't outrank twenty consistently strong ones. Volume is easy to game. Quality-adjusted ranking is not.

Screened

Every founder who enters the pipeline becomes a data point.

Funded

Which channels produce founders worth backing?

Multi-touch attribution

Credits every source that contributed to a great outcome — not just the last touch. Closes the loop back into discovery so the system compounds over time.

Met


Which channels produce founders worth meeting?

Passed

Passes are signal too — what patterns predict a pass?

How It Works — In One View

A single pipeline, an audit layer that spans it, and a feedback loop that closes it. The investment decision is derived in code — not left to the model.



The pipeline spine

- 01
- Sourcing**
Inbound decks + proactive GitHub scan against thesis.
- 02
- Screening**
Claim extraction from deck, résumé, and links.
- 03
- Diligence**
Live evidence fetch — real GitHub API, live web search.
- 04
- Decision**
Three-axis score derived in code. Audit trail attached.

What wraps it

- Agentic Audit Layer**
Spans the full pipeline. Trace log + independent validator that can only downgrade verdicts.
- Sourcing Feedback Loop**
Outcomes arc back to sourcing. The system learns which channels produce quality — and suggests where to scan next.
- Live evidence only**
Real GitHub API. Live web search. No cached data. No scripted contradictions.

Built to Be Trusted — and to Survive a Live Demo

Three design decisions that make Protegis reliable in production — and honest about what it is.



Fail-soft everywhere

A live-API outage never breaks a score. Every external dependency has a graceful fallback. The engine degrades honestly — it widens confidence bands rather than returning false precision.



No database by design

Removes the biggest deploy-day failure mode. The whole engine runs stateless. Nothing to migrate, nothing to corrupt, nothing to go stale between demo and production.



Real live data

Synthetic founders for demo purposes — but a *real* diligence engine on *real* live data. The contradiction is computed, not scripted. The GitHub repo is fetched live, every time.



The distinction matters: synthetic founders are a demo convenience. The engine itself has never seen the data it's about to check. That's what makes the ~3.2s verdict credible.

Diligence, Live. On Merit.

Every claim checked against reality. Three independent axes, never blended. The overstater, the honest unknown, and the founder who never applied — one engine.

⚡ ~3.2S · SIGNAL → DECISION

Every claim checked

Against live GitHub and live web. Not cached. Not scripted. Computed fresh, every time.

Three axes. Never blended.

Founder · Market · Idea-vs-Market. A brilliant founder in a dead market can't hide behind an average.

One engine. No exceptions.

The overstater, the legit founder, the unknown with no track record, the founder who never applied.

Capital allocated on evidence, not on who you know. That's Protegis.